

06 Linear Regression

November 20, 2022

1 Linear Regression

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from scipy import stats
```

1.1 Simple Linear Regression Model

We assume that there is approximately a linear relationship between X and Y :

$$Y \approx \beta_0 + \beta_1 * X$$

1.1.1 Least Squares Method

The least squares approach chooses $\hat{\beta}_0$ and $\hat{\beta}_1$ to minimize the residual sum of squares (RSS)

$$RSS = r_1^2 + r_2^2 + \dots + r_n^2$$

or equivalently $RSS = (y_1 - \hat{\beta}_0 - \hat{\beta}_1 x_1)^2 + \dots + (y_n - \hat{\beta}_0 - \hat{\beta}_1 x_n)^2$

1.1.2 Estimation of coefficients

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$
$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 * \bar{x}$$

where

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$$

and $\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$

1.1.3 How much do coefficient scatter?

$\hat{\beta}_0$ and $\hat{\beta}_1$ are scattered around the true values β_0 and β_1 with the standard error

$$se(\hat{\beta}_0)^2 = \sigma^2 \left(\frac{1}{n} + \frac{\bar{x}^2}{\sum_{i=1}^n (x_i - \bar{x})^2} \right)$$
$$\text{and } se(\hat{\beta}_1)^2 = \frac{\sigma^2}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

where $\sigma^2 = Var(\epsilon)$

1.1.4 Residuals and Estimation of Variance

In general σ^2 (variance of the error terms) is not known: but can be estimated on the basis of the data. The error term ϵ

- cannot be observed
- cannot be derived from $\epsilon = Y - (\hat{\beta}_0 + \hat{\beta}_1 X)$ since $\hat{\beta}_0$ and $\hat{\beta}_1$ are unknown.

Approximation for ϵ : residuals $r_i = y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i)$

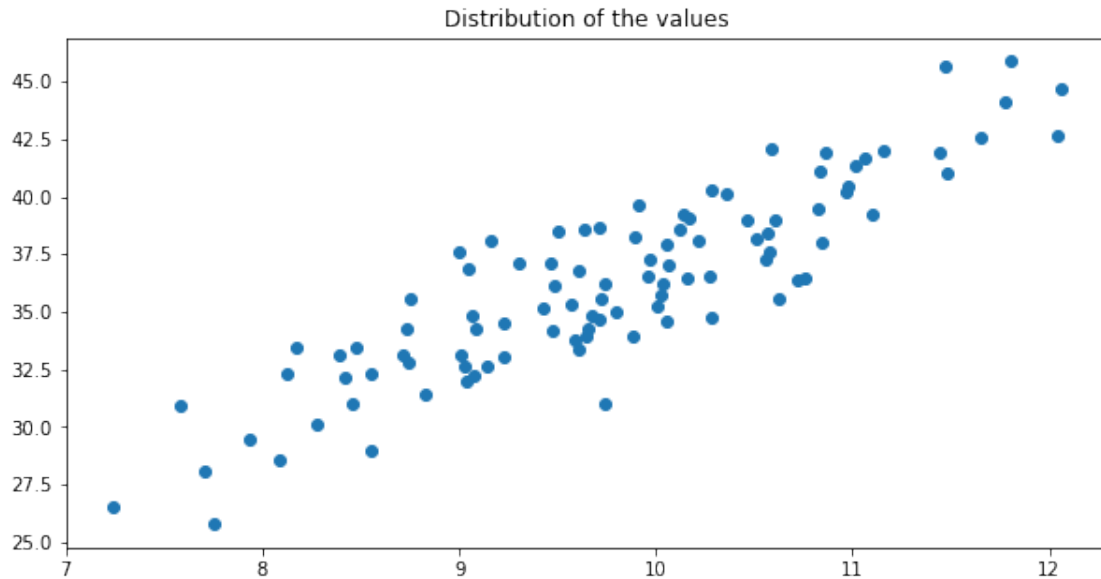
$$\hat{\sigma} = RSE = \sqrt{\frac{RSS}{n-2}} = \sqrt{\frac{r_1^2 + r_2^2 + \dots + r_n^2}{n-2}}$$

The factor $\frac{1}{n-2}$ is chosen so that the estimate of σ turns out to be unbiased. This estimate is known as the residual standard error (RSE)

```
[98]: model = stats.norm(0, 2)
      func = lambda x: 2.74 + 3.42 * x

      x = stats.norm(10, 1).rvs(100)
      y = func(x) + model.rvs(x.shape[0])
```

```
[106]: plt.figure(figsize=(10, 5))
      plt.scatter(x, y)
      plt.title("Distribution of the values")
      plt.show()
```



```
[100]: x_mean = np.mean(x)
        y_mean = np.mean(y)

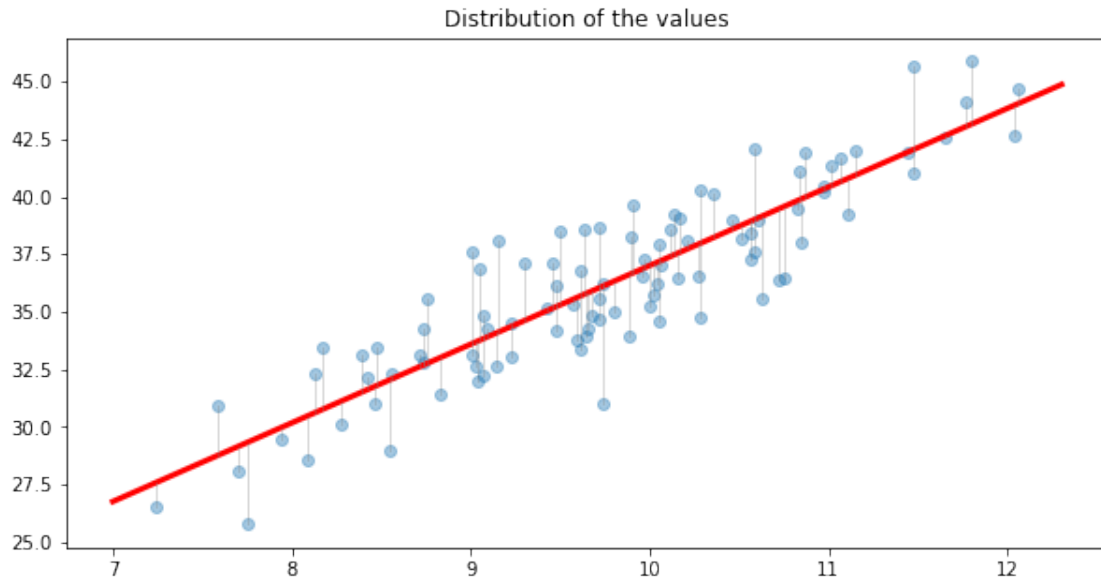
        b1 = np.sum((x - x_mean) * (y - y_mean)) / np.sum((x - x_mean) ** 2)
        b0 = y_mean - b1 * x_mean
        b1, b0
```

```
[100]: (3.4136985890489684, 2.862161280507756)
```

```
[101]: plt.figure(figsize=(10, 5))
        plt.scatter(x, y, alpha=0.4)
        plt.plot(np.array(plt.xlim()), np.array(plt.xlim()) * b1 + b0, c='r', lw=3)

        for _x, _y in zip(x, y):
            y_hat = _x * b1 + b0
            plt.plot([_x, _x], [_y, y_hat], alpha=.3, c='grey', lw=1)

        plt.title("Distribution of the values")
        plt.show()
```



```
[102]: rss = np.sum((y - (x * b1 + b0)) ** 2)
      rss
```

```
[102]: 347.577167145676
```

```
[103]: rse = np.sqrt(rss / (x.shape[0] - 2))
      rse
```

```
[103]: 1.8832699719373152
```

```
[105]: se_0 = rse ** 2 * ((1 / x.shape[0]) + (x_mean ** 2) / np.sum((x - x_mean) ** 2))
      se_1 = (rse ** 2) / np.sum((x - x_mean) ** 2)

      se_0, se_1
```

```
[105]: (3.1349680075129256, 0.03241257143945019)
```